

Memory Management

Operating-System-Concepts¹

- **paging** - a memory-management scheme allowing process's physical address space to be non-contiguous.

Linux-System-Programming²

- processes do not directly address physical memory
- each process is associated with a (unique) *virtual address space*
- address space is *linear* (start at zero) and *flat*
- **page** smallest addressable unit of memory by MMU (memory management unit)

- page size is defined by the hardware

- get pagesize via `/proc`, `getconf`, or syscall `getpagesize()`

```
# cat /proc/1/smmaps | grep -i pagesize | head -2
KernelPageSize:      4 kB
MMUPageSize:         4 kB
# getconf PAGESIZE
4096
```

```
#include <unistd.h>
```

```
#include <stdio.h>
```

```
int main() {
    printf("%d\n", getpagesize());
}
```

- pages are either valid or invalid

- * *valid page* - associated with a page of data (in RAM or secondary storage/swap)

- * *invalid page* - not associated with anything (unused, unallocated address space)

- **page fault**: - *valid* page but not currently in RAM and needs to be *paged in*

- **segmentation fault** - caused by trying to access an *invalid* page

```
#include <stdlib.h>
```

```
int main() {
    int *ip;
    ip = (int *) malloc( sizeof(int)*10 );
    ip[100000] = 99; // outside array's 10 elements, seg fault
}
```

- memory regions (areas/mappings) in every process

- *text segment* the actual program code, constants etc; ready only data

- *data segment* or *heap* - dynamic memory (`malloc()`)

- *bss segment* uninitialized global variables (zeroes)

- *stack* grows/shrinks dynamically; e.g. local vars, function return data. Thread: one stack per thread

```
$ ulimit -Sa | grep -i stack
```

```
stack size          (kbytes, -s) 8192
```

¹<https://codex.cs.yale.edu/avi/os-book/>

²<https://www.oreilly.com/library/view/linux-system-programming/9781449341527/>

- **dynamic memory** - allocated at runtime
 - `void * malloc (size_t size);`
 - `calloc()` - zero's out bytes ('c'=clear)
 - `realloc()` > if unable to enlarge the existing chunk of memory by growing the chunk in situ, > the function may allocate a new region of memory size bytes in length, > copy the old region into the new one, and free the old region
- **anonymous memory mapping**
 - for large allocations, glibc doesn't use the heap; instead creates an anonymous memory mapping
 - `mmap()` creates a memory mapping and the system call `munmap()` destroys a mapping
- **advanced memory allocation**
 - `mallopt()` - sets the memory-management-related parameter e.g.
 - * `M_MMAP_THRESHOLD` - malloc allocation request via an anonymous mapping instead of heap
 - * default value (Linux) if not set for `M_MMAP_THRESHOLD` = 128k (128*2024)
- **stack-based allocations**
 - one can use the stack for dynamic memory allocations too (as long as it doesn't overflow)
 - `alloca()` system call for dynamic memory allocation from stack > `alloca()` returns a pointer to size bytes of memory. > This memory lives on the stack and is automatically freed when the invoking function returns.

```
#include <alloca.h>
#include <stdlib.h>
int main() {
    void *p;
    // int size = 8388608; /* will segfault as ulimit: stack size (kbytes, -s) 8192 */
    int size = 8388608 / 2;
    p = alloca(size);
    // free(p); /* will result in Aborted (core dumped) */
}
```

 - very fast compared to `malloc()` (just increases stack pointer)
- **locking memory** - two situations you might want to change default paging behaviour:
 - (1) *determinism* - page fault are costly b/c of I/O operations (disk)
 - (2) *security* - if secrets kept in memory are paged out to (unencrypted) disk
 - `mlock()` - to locks a specific memory area/allocation, unlock with `munlock()`
 - `mlockall()` - locks the whole process space, unlock with `munlockall()`